


THE DATA BOTTLENECK



Why offline-first infrastructure matters for reproducible psychiatry


SYSTEM STACK

Zero data egress architecture for clinical audio intelligence



THE INTERFACE

Annotated clinical intelligence view



Sentiment Score

Risk Flag Proxy


Speaker-Tagged Transcript

THE TERMINAL

Open Source & Reproducible

pipeline.py

```
1 import whisper
2 from pyannote.audio import Pipeline
3
4 def run_pipeline(audio_path: str):
5     model = whisper.load_model("base")
6     result = model.transcribe(audio_path)
7     dia = Pipeline.from_pretrained("pyannote/speaker-diarization")
8     segments = dia(audio_path)
9     print("Sentiment score:", result["sentiment"])
10    return result, segments
11
12 # Built on OpenAI Whisper + Pyannote + Python
13 if __name__ == "__main__":
14     run_pipeline("session.m4a")
```



QR PLACEHOLDER

DEPLOY THE INFRASTRUCTURE.

huggingface.co/spaces/ChengdongPeter/Clinical-Whisper
github.com/czhou732/Clinical-Whisper-Pipeline

Chengdong (Peter) Zhou | USC