

ClinicalWhisper 5.2: getting started

Installing

You'll need a Mac with an M1 chip or newer, macOS 14 or later, and about 25 GB of free space. Scoring works best with 16 GB of memory or more; on smaller Macs the app starts with scoring turned off.

Drag ClinicalWhisper into Applications. The app isn't registered with Apple yet, so macOS will block it the first time. On macOS 14, right-click the app and choose Open. On macOS 15 or newer, try opening it once, then go to System Settings > Privacy & Security, scroll down, and click Open Anyway. If neither works, paste this into Terminal:

```
xattr -dr com.apple.quarantine /Applications/ClinicalWhisper.app
```

The first launch takes about a minute. Before running real recordings, please make sure FileVault is on (System Settings > Privacy & Security), since transcripts are saved to the disk.

Running recordings

Drop files onto the window or click to pick them. Most audio formats work (wav, mp3, m4a, flac and so on). The app reads them where they are and doesn't move or copy them.

Set "People in the recording" to the number of speakers. For a one-on-one interview, leave it at 2. Getting this right matters more than anything else for the speaker labels.

If you only need transcripts, tick "Transcribe only". It skips the clinical scores and is a lot faster. Names and other identifiers are masked either way.

You'll see a time estimate before you start. Keep the laptop plugged in with the lid open, and you'll get a notification when it's done. If the Interviewer label lands on the wrong person, pick the right one from the Interviewer menu and click Apply.

Results

Everything is saved in a folder called ClinicalWhisper in your home folder. It's deliberately not in Documents, so iCloud won't upload it. "Open results folder" in the app takes you there. Each recording gets an `_analysis.json` file with the masked transcript, voice features and scores, and each batch gets a summary CSV.

Names and other identifiers are replaced with tags like `[first_name_1]`. The same tag always means the same word, and there's a key at the top of each transcript. The masking is automatic and it does miss things sometimes, so please read over a transcript before sharing it with anyone.

Each result also reports how far the voice sits above the background noise (SNR). Under 15 dB it's flagged, and the voice measures from that recording shouldn't be trusted. Response latency isn't reliable in this version yet, so please leave it out of any analysis.

Lots of files

For a big folder, Terminal is easier:

```
cd /Applications/ClinicalWhisper.app/Contents/MacOS
./ClinicalWhisper --batch -i "/path/to/recordings" -o ~/ClinicalWhisper/summary.csv \
  --speakers 2 --transcribe-only
```

It prints progress and the time left as it goes. If it gets interrupted, run the same command again with `--resume` added and it picks up where it stopped. Add `--ids ids.csv` to attach participant IDs (columns filename, participant_id, session_label), and `--help` lists the rest. Try not to run a batch while the app window is also processing, since they'll fight over memory.

For a sense of timing: on my M2 Max, a three-and-a-half-hour interview took about 12 minutes to transcribe and another 17 to score. Yours may be slower. The estimate adjusts after a few files.

If something breaks

Click "Save diagnostics for a bug report" at the bottom of the window and email me the zip. It has the logs and some basic info about the Mac, with file names stripped out, and no audio or transcripts.

One last thing: this is a research tool, not a medical device, and the scores aren't meant for diagnosing anyone.

Peter